

Pet-Match Recommender – Project Outline

1 Project Overview

The goal is to build a **pet-adopter matching prototype** that shows how generative AI and advanced analytics can raise shelter adoption rates. In four weeks our deliverables will be:

- A Streamlit demo that recommends the top-3 pets for any adopter persona.
- Report detailing the business plan, analytics plan, and effect of the tool

The trick is to treat your synthetic pets table as **item data**, invent a small but consistent **adopter-profile table** plus **positive pairs**, and then train a *content-based* or *metric-learning* recommender.

You can layer an LLM on top to generate friendly “why this pet?” explanations.

1 Create an adopter table and pseudo-labels

Step	How to do it	Lines of code
1 a. Build 5–10 adopter persona s	Use faker to sample age, housing, activity, prior-pet; save 5 000–10 000 rows.	<pre>python\nfrom faker import Faker\nfaker =\nFaker()\nadopters = [{\n 'adopter_id':\n faker.uuid4(),\n 'age':\n faker.random_int(21,70),\n 'housing':\n faker.random_element(['apt', 'house', 'farm']),\n 'activity':\n faker.random_element(['low', 'mod', 'high']),\n 'has_prior_pet': faker.boolean()\n}] for _ in\nrange(8000)]</pre>

1 b.	<i>If housing=apt ⇒ prefer “Small” breed; if activity=high ⇒ prefer “Working/Active ” breeds, etc.</i>	Write a Spark UDF that scores (adopter, pet) pairs on 0–1.
1 c.	Cross-join adopters × ~1 000 random pets, keep rows where rule-score > 0.7 and label = 1; sample same # for label = 0.	Spark: <code>adopters.crossJoin(pets_silver.sample(0.05)).withColumn("label", udf_score(...))</code>

You now have **tens of thousands of labelled (adopter, pet) pairs** for supervised training.

2 Model options

Model	Why it works label-light	Tooling
Wide-&-Deep tabular	Handles one-hot (breed, color) + numeric (age). Learns interactions without huge data.	TensorFlow Keras <code>(tf.estimator.DNNLinearCombinedClassifier)</code> .
Siamese / Triplet net	Learns embeddings for pets & adopters; inference = cosine similarity.	PyTorch, 100 k pairs fit in Colab GPU.

LightFM hybrid	Uses content features + (synthetic) interactions; easy to evaluate Precision@k.	<code>lightfm</code> Python package.
-----------------------	---	--------------------------------------

Pick one; Wide-&-Deep is easiest.

3 Streamlit demo flow

adopter dropdown → `model.predictTopK(pet_vecs)` →

LLM tells the story → user clicks 👍 / 👎 → write new feedback row

LLM explanation

```
explain_prompt = (
```

```
    f"You are a pet-placement counselor.\n"
```

```
    f"Adopter profile: {adopter_json}\n"
```

```
    f"Pet profile: {pet_json}\n"
```

```
    f"Explain in one short paragraph why this is a good match.")
```

```
response = openai.ChatCompletion.create(model="gpt-3.5-turbo", messages=[...])
```

Cache the explanation along with the recommendation.

4 Putting it all in Spark + Streamlit

1. `01_make_adopters.ipynb`

Generate adopter table, write `adopters_silver` to shared drive.

2. 02_make_pairs.ipynb

Cross-join, score rule, create *train_pairs.parquet* (features + label).

3. 03_train_wd_model.ipynb

Spark → Pandas sample → TensorFlow training → save *pet_model.h5*.

app.py (Streamlit)

```
import tensorflow as tf, pandas as pd, streamlit as st

model = tf.keras.models.load_model("pet_model.h5")

pets = pd.read_parquet("pets_silver")      # small subset in memory

adopter = st.selectbox("Choose adopter persona", adopters_table["adopter_id"])

topk = rank_pets(model, adopter, pets, k=3)

for pet in topk: st.image(pet["photo"]); st.write(explain(pet, adopter))
```

4.

5 Why graders will accept synthetic labels

- You document the rule-based generator and label it “**synthetic for prototyping only.**”
- The Streamlit demo is interactive and explainable through the LLM.

That combination ticks the “creative + effective code/tool” rubric